

PHASE 1

once per judge / task type

Calibrate

50 pairs $\times$ 3–5 variants

Gate

any bias
 > 10 pp?

yes

no

Reject Judge

(exit 1)

Unstable

promotion (exit)

pp = percentage points

PHASE 2

per model version

Evaluate

$k = 3$ variants
($3 \times$ cost)

Gate

range
 > 15 pp?

yes

no

Promote

gate passes (exit)

Promote